


Week 11

Start at 4:05

P1: Interleavings of Threads

Critical section: Two or more code-parts that access and manipulate shared data (aka a shared resource).

Race condition: “a situation in which multiple threads read and write a shared data item and the final result depends on the relative timing of their execution”.



P2: Lock-based Thread Synchronisation

Mutex

Example

Blocking

```
pthread_mutex_t mylock = PTHREAD_MUTEX_INITIALIZER;

void * thread_function(void * arg) {
    /*
        Assume There are 3 threads.

        T1 locks and executing critial section
        T2, T3 are waiting in the line `pthread_mutex_lock(&mylock);`

        Do not know after unlock mutex which thread will execute (order)
    */

    pthread_mutex_lock(&mylock);
    counter = counter + 1; //critical section
    pthread_mutex_unlock(&mylock);
}
```

Non-blocking

```
pthread_mutex_t mylock = PTHREAD_MUTEX_INITIALIZER;

void * thread_function(void * arg) {
    /*
        Assume There are 3 threads.

        T1 try lock successfully and this function locks automatically
        and executing critial section
        T2, T3 try lock fail and printf("Unscu...."), continue executing
    */
    if ( 0 != pthread_mutex_trylock(&mylock) ) {
        printf("Unsuccessful attempt to acquire lock\n");
    }
}
```

```

        // ##### pthread_mutex_unlock(&mylock); // Error: mutex not
        acquired! ==> Only the thread that owns a mutex should unlock it!
        #####

        // what happens => unlock thread2

    } else {
        // critical section start:
        //... ==> thread1 still goes here => can cause race
        condition as well
        // critical section end
        pthread_mutex_unlock(&mylock);
    }
}

```

Dynamic creation of mutexes

```

pthread_mutex_t * mylock;

mylock = (pthread_mutex_t *) malloc(sizeof(pthread_mutex_t));
pthread_mutex_init(mylock, NULL);

...

pthread_mutex_destroy(mylock);
free (mylock);

```

Serialization

```

/*
    Part 1, Serialization: one thread has to wait another
*/
unsigned long counter=0;

```

```

void * thread_function(void * arg) {
    long l;
    for (l=0; l < MAX_ITER; l++) {
        pthread_mutex_lock(&mylock);
        counter = counter + LongComputation();

        pthread_mutex_unlock(&mylock);
    }
}

/*
    Part 2, Reduce Serialization
*/
unsigned long counter=0;

void * thread_function(void * arg) {
    long l, tmp;
    for (l=0; l < MAX_ITER; l++) {
        tmp = LongComputation();

        pthread_mutex_lock(&mylock);
        counter=counter+tmp;//crit.sect.
        pthread_mutex_unlock(&mylock);
    }
}

```

P3: DeadLock

Example

T1

T2

cat 1

fox 1

dog 2

dog 3 (blocking)

fox (blocking)

Can not reach next, can not continue to unlock

free...

free...

Necessary conditions for a deadlock

1. Mutual exclusion: a resource can be assigned to at most one thread.
2. Hold and wait: threads both hold resources and request other resources.
3. No preemption: a resource can only be released by the thread that holds it.
4. Circular wait: a cycle exists in which each thread waits for a resource that is assigned to another thread.

Using 10 minutes to write Q2, back at 4:50

DO not refer to Lecture answer, write by yourself

Deadlock Prevention

Cycles can be prevented by a locking hierarchy: ==> breaking rule 4

1. Impose an ordering on mutexes.
2. Require that all threads acquire mutexes in the same order.

==> a special case

Thread 0:

```
acquire mutex A
for(;;) // enter endless loop
    // never free mutex A
```

Thread 1:


```
try to acquire mutex A
wait (block) forever for mutex A
```



Example: Dining Philosopher's Deadlock (Skip this)

```
#define MAX 5
pthread_t thr[MAX];
pthread_mutex_t m[MAX];

void * tfunc (void * arg) {
    long i = (long) arg; // thread id: 0..4
    for (;;) {
        pthread_mutex_lock( &m[i] );
        pthread_mutex_lock( &m[(i + 1) % MAX] );
        printf("Philosopher %d is eating...\n", i);
        pthread_mutex_unlock(&m[i]);
        pthread_mutex_unlock(&m[(i + 1) % MAX]);
    }
}
```




P0,...P3: m[0]→m[1]→m[2]→m[3]→m[4] 😊

P4: m[4]→m[0] 😞

Consider the case for MAX=2:

```
acquire chopstick 0
try to acquire chopstick 1
wait for chopstick 1
```




```
acquire chopstick 1
try to acquire chopstick 0
wait for chopstick 0
```

Dining Philosopher's (fixed)

- Introduce a locking hierarchy:
 - pick up the chopstick with the smaller index.
 - Pick up the chopstick with the higher index.

```
for (;;) {
    if ( i < ((i + 1) % MAX) ) {
        pthread_mutex_lock(&mtx[i]);
        pthread_mutex_lock(&mtx[(i + 1) % MAX]);
    } else {
        pthread_mutex_lock(&mtx[(i + 1) % MAX]);
        pthread_mutex_lock(&mtx[i]);
    }
    printf("Philosopher %d is eating...\n", i);
    pthread_mutex_unlock(&mtx[i]);
    pthread_mutex_unlock(&mtx[(i + 1) % MAX]);
}
```



locking hierarchy:
P0,...P4: m[0]→m[1]→m[2]→m[3]→m[4] 😊

Consider the case for MAX=2:

```
acquire chopstick 0
acquire chopstick 1
eat
release chopstick 0
release chopstick 1
```



```
try to acquire chopstick 0
wait for chopstick 0
...
...
...
```

41

Hold and wait ==> breaking rule 2

T1

cat

operation(cat)

dog

unlock cat

lock(dog)

operation(dog)

fox

unlock dog

lock(cat, fox)


operation(cat, fox)


Starvation and Live-lock (Extension)

Starvation happens when “greedy” threads make shared resources unavailable for long periods. For instance, suppose an object provides a synchronized method that often takes a long time to return. If one thread invokes this method frequently, other threads that also need frequent synchronized access to the same object will often be blocked.

P4: Lock Contention and Scalability

- Coarse-grained lock => a lock for the whole linked list => no parallel => we can not do search in parallel
- Medium-grained lock.
- **Fine-grained lock.**





```
lock A ?  
--- No  
  
lock B ?  
--- No  
  
lock C ?  
--- Yes      ==> the other thread can not access nodes after C  
              ==> some race condition may happens when two threads  
both inserting  
              after C without lock  
  
lock D as well  
update(C, D)  
  
unlock C, D
```

P5: Semaphores

Semaphores are non-negative integer synchronization variables.

If you want to do something with an order.

Something like wake up using signal are without order

s has some init value maybe 0

V(s): [s++;] ==> sem_post

```
P(s):  
[  
    while (s == 0) {  
        wait();  
    }  
    s--;  
] ==> sem_wait
```

The statements between brackets [] are therefore an atomic operation.

==> At any time, only one P() or V() operation can modify s.

Example:

```
#include <semaphore.h>  
#define MAX 4  
#define MAX_ITER 50000000  
  
pthread_t thr[MAX];  
sem_t s;  
long counter = 0;  
  
void * tfunc (void * arg) {  
    int i;  
    for (i=0; i<MAX_ITER; i++) {  
        sem_wait(&s);  
        counter++; //critical section  
        sem_post(&s);  
    }  
}
```

```

}
}
int main() {
    int i, j;
    sem_init(&s, 0, 1);
    ...
}

```

Example: Dining Philosophers

```

sem_t chpstcks[N], limit;

void * Philosopher(void * arg) {
    long id = (long) arg;
    for(;;) {
        think();
        sem_wait(&limit);
        sem_wait(&chpstcks[id]);
        sem_wait(&chpstcks[(id+1)%N]);
        eat();
        sem_post(&chpstcks[id]);
        sem_post(&chpstcks[(id+1)%N]);
        sem_post(&limit);
    }
}

int main() {
    int i;
    for(i=0; i<N; i++)
        sem_init(&chpstcks[i], 0, 1);
    sem_init(&limit, 0, N-1);
    ...
}

```

- A counting semaphore can prevent the Dining Philosophers from dead-locking:
 - Assume N philosophers sitting at the table.
 - Use a counting semaphore with an initial count of N-1.
 - At most N-1 philosophers can pick up the left chopstick at once.
 - At least one of those philosophers will have access to two chopsticks. This philosopher can eat.



Synchronizing Threads using Semaphores

```
sem_t s;

void * T1(void * arg) {
    ...
    printf("this comes first\n");
    sem_post(&s);
    ...
}

void * T2(void * arg) {
    ...
    sem_wait(&s); //wait for T1
    printf("this comes second\n");
    ...
}

int main() {
    sem_init(&s, 0, 0);
    ...
}
```

- Besides mutual exclusion, semaphores can also be used to synchronize threads.
- Example:
 - Assume 2 threads executing the thread routines T1 and T2.
 - Assume a semaphore s.
 - s is initialized to 0.
 - T2 has a sem_wait() operation on s.
 - T1 has a sem_post() operation.
 - When T2 reaches sem_wait(), it will block until T1 has executed sem_post().
 - Question: what happens if T1 executes sem_post() before T2 executes sem_wait() ?

57

Question: what happens if T1 executes sem_post() before T2 executes sem_wait() ?

Synchronizing Threads using Semaphores

```
1 sem_t _____;  
2  
3 void * T1(void * arg) {  
4     for (;;) {  
5         _____  
6         printf ("ping\n");  
7         _____  
8     }  
9 }  
10  
11 void * T2(void * arg) {  
12     for (;;) {  
13         _____  
14         printf ("pong\n");  
15         _____  
16     }  
17 }  
18 int main() {  
19     _____  
20     _____  
21     ...  
22 }
```

- Example:

- Assume 2 threads, executing the thread routines T1 and T2, respectively.
- Thread T1 outputs "ping" in an endless loop.
- Thread T2 outputs "pong" in an endless loop.
- How can we synchronize T1 and T2 using semaphores, such that the output will be

ping
pong
ping
pong
...

?

58

Week11/Ping-Pong , Q3 & Q4

Before Q4, Intro Barrier

10 minutes Break => continue Q3 back 5:23

We have previously solved the dining philosophers problem by using a locking hierarchy.

This time, use a semaphore for the table that only allows $N/2$ philosophers to eat at a time

DO Question 4 ==>

Come up with an idea

10 min back at 5:50

P6: Amdahl's Law

- p = fraction of work that can be parallelized.
- n = the number of threads executing in parallel.

$$\text{Speedup} = \frac{\text{old_running_time}}{\text{new_running_time}} = \frac{1}{(1 - p) + \frac{p}{n}}$$

$(1 - p)$ => the part can not be parallelised

P => the part can be parallelised, and we have n workers.

If 100% => $p = 1$ => speed = n (linear)

Amdahl's Law

